

# Linux Device Drivers

- Deferred Works



# Agenda

- Concept of jiffies
- Kernel Timers
- Tasklets
- Work Queues
- Shared Queue



# Concept of jiffies



# Concept of Jiffies

- Timer interrupts are generated by the system's timing hardware at regular intervals
- The interval, as per the Linux kernel is given by '*HZ*' ticks per second. For example, if the value of HZ is 250(say), that means that the timer interrupts 250 times per second.
- Scheduling and other timely functionalities depends on this value of HZ.
- Each time a timer interrupt occurs, the value of an internal 64-bit(even on 32-bit platforms) counter, named *jiffies\_64* is incremented.
- Driver authors can normally access *jiffies* which is an *unsigned long* of the least 32-bits of *jiffies\_64*.

# Applying jiffies : Profiling

- unsigned long t1, t2, diff, msec;

...

t1 = jiffies;

/\* Code you wish to profile \*/

t2 = jiffies;

diff = t2 - t1;

msec = diff \* 1000 / HZ;

# Applying jiffies : Generating delays

- The following macros are used to work with jiffies
- `<linux/jiffies.h>`

`int time_after(unsigned long a, unsigned long b);`

`int time_before(unsigned long a, unsigned long b);`

`int time_after_eq(unsigned long a, unsigned long b);`

`int time_before_eq(unsigned long a, unsigned long b);`

# Applying jiffies : Generating delays

- General delay calculation :  
    `unsigned long d = jiffies + MILLI_DELAY * HZ / 1000;`
- `while (time_before(jiffies, d))`  
    `schedule();`

# Short delays

- Busy waiting
  - `void ndelay(unsigned long nsecs);`
  - `void udelay(unsigned long usecs);`
  - `void mdelay(unsigned long msecs);`
- Delayed sleeping
  - `void msleep(unsigned long msecs);`
  - `void ssleep(unsigned long secs);`



# Kernel Timers

# Kernel Timers

- A kernel timer is a data structure that instructs the kernel to execute a user-defined function with a user-defined argument at a user defined time.
- Kernel timers are run as the result of a “software interrupt” i.e. they be executed atomically, without any sleeps involved.
- The timer API :
  - `<linux/timer.h>`

```
struct timer_list {  
    /* ... */  
    unsigned long expires;  
    void (*function)(unsigned long);  
};
```

# Kernel Timers : Usage

- `/* Initialization */`

```
void timer_setup(struct timer_list *timer,void (*fun)(struct timer_list*),flags); // refer header-file for flags.
```

- `/* Timer Operations */`

```
void add_timer(struct timer_list *timer);
```

```
int mod_timer(struct timer_list *timer, unsigned long expires);
```

- `/* Deletion */`

```
int del_timer_sync(struct timer_list *timer);
```



# Tasklets



# Tasklets

- Tasklets resemble kernel timers in some ways. Unlike kernel timers, however, you can't ask to execute the function at a specific time.
- By scheduling a tasklet, you simply ask for it to be executed at a later time chosen by the kernel.
- This behavior is especially useful with interrupt handlers, where the hardware interrupt must be managed as quickly as possible, but most of the *data management can be safely delayed to a later time*.
- Tasklet API :
  - `<linux/interrupt.h>`

```
struct tasklet_struct {  
    /* ... */  
    void (*function)(unsigned long);  
    unsigned long data;  
};
```

# Tasklets : Usage

- `/* Initialization */`  
`tasklet_init(struct tasklet_struct *t, void (*func)(unsigned long),`  
`unsigned long data);`
- `/* Operations on tasklets */`  
`tasklet_schedule(struct tasklet *t);`  
`tasklet_disable(struct tasklet *t);`  
`tasklet_enable(struct tasklet *t);`
- `/* Destroying the tasklet */`  
`tasklet_kill(struct tasklet *t);`



# Work Queues



# Work Queues

- *Workqueues* are, superficially, similar to tasklets, but with a difference that workqueues run in process context.
- Usage :

`<linux/workqueue.h>`

```
struct workqueue_struct;
```

```
struct work_struct;
```



# Work Queues : Usage

- /\* Work queue creation \*/  
struct workqueue\_struct \*create\_singlethreaded\_workqueue(char \*name);
- /\* Create a work \*/  
INIT\_WORK(struct work\_struct \*work, void (\*func)(struct work \*));
- /\* Assign the work to the work queue \*/  
int queue\_work(struct workqueue\_struct \*queue,  
                struct work\_struct \*work);
- /\* Operations on work queues \*/  
void flush\_workqueue(struct workqueue \*queue);  
void destroy\_workqueue(struct workqueue \*queue);

# Shared Queue

- It is similar to that of a work queue except that the queue is maintained by the kernel itself.
- Usage :
  - `<linux/workqueue.h>`  
`struct work_struct;`
  - `/* Initialise the work */`  
`INIT_WORK(struct work_struct *work, void (*func)(struct work *));`
  - `/* Schedule the work to execute */`  
`int schedule_work(struct work_struct *work);`

# References

- Jonathan Corbet, Alessandro Rubini and Greg Kroah-Hartman, "Linux Device Drivers", 3rd Edition, O'Reilly Publications
- Robert Love, "Linux Kernel Development", 3rd Edition, Developer's Library

*Thank You :)*