

Looking at the Sources

Kernel boot starts with the execution of real mode assembly code living in the *arch/x86/boot/* directory. Look at *arch/x86/kernel/setup_32.c* to see how the protected mode kernel obtains information gleaned by the real mode kernel.

The first boot message is printed by code residing in *init/main.c*. Dig inside *init/calibrate.c* to learn more about BogomIPS calibration and *include/asm-your-arch/bugs.h* for an insight into architecture-dependent checks.

Timer services in the kernel consist of architecture-dependent portions that live in *arch/your-arch/kernel/* and generic portions implemented in *kernel/timer.c*. For related definitions, look at the header files, *include/linux/time*.h*.

jiffies is defined in *linux/jiffies.h*. The value for HZ is processor-dependent and can be found in *include/asm-your-arch/param.h*.

Memory management sources reside in the top-level *mm/* directory.

Table 2.1 contains a summary of the main data structures used in this chapter and the location of their definitions in the source tree. Table 2.2 lists the main kernel programming interfaces that you used in this chapter along with the location of their definitions.

TABLE 2.1 Summary of Data Structures

Data Structure	Location	Description
HZ	<i>include/asm-your-arch/param.h</i>	Number of times the system timer ticks in 1 second
loops_per_jiffy	<i>init/main.c</i>	Number of times the processor executes an internal delay-loop in 1 jiffy
timer_list	<i>include/linux/timer.h</i>	Used to hold the address of a routine that you want to execute at some point in the future
timeval	<i>include/linux/time.h</i>	Timestamp
spinlock_t	<i>include/linux/spinlock_types.h</i>	A busy-locking mechanism to ensure that only a single thread enters a critical section
semaphore	<i>include/asm-your-arch/semaphore.h</i>	A sleep-locking mechanism that allows a predetermined number of users to enter a critical section
mutex	<i>include/linux/mutex.h</i>	The new interface that replaces semaphore
rwlock_t	<i>include/linux/spinlock_types.h</i>	Reader-writer spinlock
page	<i>include/linux/mm_types.h</i>	Kernel's representation of a physical memory page

TABLE 2.2 Summary of Kernel Programming Interfaces

Kernel Interface	Location	Description
time_after() time_after_eq() time_before() time_before_eq()	<i>include/linux/jiffies.h</i>	Compares the current value of <i>jiffies</i> with a specified future value
schedule_timeout()	<i>kernel/timer.c</i>	Schedules a process to run after a specified timeout has elapsed
wait_event_timeout()	<i>include/linux/wait.h</i>	Resumes execution if a specified condition becomes true or if a timeout occurs
DEFINE_TIMER()	<i>include/linux/timer.h</i>	Statically defines a timer
init_timer()	<i>kernel/timer.c</i>	Dynamically defines a timer
add_timer()	<i>include/linux/timer.h</i>	Schedules the timer for execution after the timeout has elapsed
mod_timer()	<i>kernel/timer.c</i>	Changes timer expiration
timer_pending()	<i>include/linux/timer.h</i>	Checks if a timer is pending at the moment
udelay()	<i>include/asm-your-arch/delay.h</i> <i>arch/your-arch/lib/delay.c</i>	Busy-waits for the specified number of microseconds
rdtsc()	<i>include/asm-x86/msr.h</i>	Gets the value of the TSC on Pentium-compatible processors
do_gettimeofday()	<i>kernel/time.c</i>	Obtains wall time
local_irq_disable()	<i>include/asm-your-arch/system.h</i>	Disables interrupts on the local CPU
local_irq_enable()	<i>include/asm-your-arch/system.h</i>	Enables interrupts on the local CPU
local_irq_save()	<i>include/asm-your-arch/system.h</i>	Saves interrupt state and disables interrupts
local_irq_restore()	<i>include/asm-your-arch/system.h</i>	Restores interrupt state to what it was when the matching <i>local_irq_save()</i> was called
spin_lock()	<i>include/linux/spinlock.h</i> <i>kernel/spinlock.c</i>	Acquires a spinlock.
spin_unlock()	<i>include/linux/spinlock.h</i>	Releases a spinlock
spin_lock_irqsave()	<i>include/linux/spinlock.h</i> <i>kernel/spinlock.c</i>	Saves interrupt state, disables interrupts and preemption on local CPU, and locks their critical section to regulate access by other CPUs
spin_unlock_irqrestore()	<i>include/linux/spinlock.h</i> <i>kernel/spinlock.c</i>	Restores interrupt state and preemption and releases the lock
DEFINE_MUTEX()	<i>include/linux/mutex.h</i>	Statically declares a mutex

Continues

TABLE 2.2 Continued

Kernel Interface	Location	Description
<code>mutex_init()</code>	<i>include/linux/mutex.h</i>	Dynamically declares a mutex
<code>mutex_lock()</code>	<i>kernel/mutex.c</i>	Acquires a mutex
<code>mutex_unlock()</code>	<i>kernel/mutex.c</i>	Releases a mutex
<code>DECLARE_MUTEX()</code>	<i>include/asm-your-arch/semaphore.h</i>	Statically declares a semaphore
<code>init_MUTEX()</code>	<i>include/asm-your-arch/semaphore.h</i>	Dynamically declares a semaphore
<code>up()</code>	<i>arch/your-arch/kernel/semaphore.c</i>	Acquires a semaphore
<code>down()</code>	<i>arch/your-arch/kernel/semaphore.c</i>	Releases a semaphore
<code>atomic_inc()</code> <code>atomic_inc_and_test()</code> <code>atomic_dec()</code> <code>atomic_dec_and_test()</code> <code>clear_bit()</code> <code>set_bit()</code> <code>test_bit()</code> <code>test_and_set_bit()</code>	<i>include/asm-your-arch/atomic.h</i>	Atomic operators to perform light-weight operations
<code>read_lock()</code> <code>read_unlock()</code> <code>read_lock_irqsave()</code> <code>read_lock_irqrestore()</code> <code>write_lock()</code> <code>write_unlock()</code> <code>write_lock_irqsave()</code> <code>write_lock_irqrestore()</code>	<i>include/linux/spinlock.h</i> <i>kernel/spinlock.c</i>	Reader-writer variant of spinlocks
<code>down_read()</code> <code>up_read()</code> <code>down_write()</code> <code>up_write()</code>	<i>kernel/rwsem.c</i>	Reader-writer variant of semaphores
<code>read_seqbegin()</code> <code>read_seqretry()</code> <code>write_seqlock()</code> <code>write_sequnlock()</code>	<i>include/linux/seqlock.h</i>	Seqlock operations
<code>kmalloc()</code>	<i>include/linux/slab.h</i> <i>mm/slab.c</i>	Allocates physically contiguous memory from <code>ZONE_NORMAL</code>
<code>kzalloc()</code>	<i>include/linux/slab.h</i> <i>mm/util.c</i>	Obtains zeroed kmalloced memory
<code>kfree()</code>	<i>mm/slab.c</i>	Releases kmalloced memory
<code>vmalloc()</code>	<i>mm/vmalloc.c</i>	Allocates virtually contiguous memory that is not guaranteed to be physically contiguous.